


# Solution to COMP9318 Assignment 1

## Q1

1. See the figure below.



Note:

- In STATUS, it is possible to replace **num\_year** with **start\_date**. The current design is easy to query but need continuous update.
- The current design models each course offerings in COURSE. If we choose to model each course (by removing SESSION, YEAR, and LECTURER), then we need to have another dimension table named LECTURER.

2. The tables are shown below.

| STUDENT           |                   |                  |                  |                 |                |                | STATUS           |                 |                |             |
|-------------------|-------------------|------------------|------------------|-----------------|----------------|----------------|------------------|-----------------|----------------|-------------|
| <i>studentKey</i> | <i>sid</i>        | <i>sname</i>     | <i>suburb</i>    | <i>postcode</i> | <i>dept</i>    | <i>faculty</i> | <i>statusKey</i> | <i>num_year</i> |                |             |
| 1                 | 111               | John Doe         | Kingsford        | 2033            | CSE            | Engineering    | 1                | 3               |                |             |
| 2                 | 112               | John Smith       | Kingsford        | 2033            | EE             | Engineering    | 2                | 3               |                |             |
| COURSE            |                   |                  |                  |                 |                |                | TIME             |                 |                |             |
| <i>courseKey</i>  | <i>cid</i>        | <i>cname</i>     | <i>lecturer</i>  | <i>year</i>     | <i>session</i> | <i>dept</i>    | <i>faculty</i>   | <i>timeKey</i>  | <i>session</i> | <i>year</i> |
| 1                 | 101               | COMP6714         | Wei Wang         | 2015            | 2              | CSE            | Engineering      | 1               | 2              | 2015        |
| 2                 | 102               | COMP6714         | Wei Wang         | 2016            | 2              | CSE            | Engineering      | 2               | 2              | 2016        |
| 3                 | 103               | COMP9318         | Wei Wang         | 2016            | 2              | CSE            | Engineering      | 3               | 2              | 2016        |
| ENROLLMENT        |                   |                  |                  |                 |                |                |                  |                 |                |             |
|                   | <i>studentKey</i> | <i>statusKey</i> | <i>courseKey</i> | <i>timeKey</i>  | <i>mark</i>    |                |                  |                 |                |             |
|                   | 1                 | 1                |                  | 1               | 71             |                |                  |                 |                |             |
|                   | 1                 | 1                |                  | 3               | 83             |                |                  |                 |                |             |
|                   | 2                 | 2                |                  | 3               | 86             |                |                  |                 |                |             |
|                   | 2                 | 2                |                  | 2               | 69             |                |                  |                 |                |             |

### 3. MDX Query

```
SELECT {[Student].[Department].MEMBERS} on COLUMNS
      {[Time].[Year].MEMBERS}
FROM [Enrollment]
WHERE (Measures.[mark], [Course].[Faculty].[Engineering])
```

## Q2

1. See the following table.

| <i>cuboid</i> | <i>Location</i> | <i>Time</i> | <i>Item</i> | <i>SUM(Quantity)</i> |
|---------------|-----------------|-------------|-------------|----------------------|
| LTI           | Sydney          | 2005        | PS2         | 1400                 |
| LTI           | Sydney          | 2006        | PS2         | 1500                 |
| LTI           | Sydney          | 2006        | Wii         | 500                  |
| LTI           | Melbourne       | 2005        | XBox 360    | 1700                 |
| LT            | Sydney          | 2005        | ALL         | 1400                 |
| LT            | Sydney          | 2006        | ALL         | 2000                 |
| LT            | Melbourne       | 2005        | ALL         | 1700                 |
| LI            | Sydney          | ALL         | PS2         | 2900                 |
| LI            | Sydney          | ALL         | Wii         | 500                  |
| LI            | Melbourne       | ALL         | XBox 360    | 1700                 |
| TI            | ALL             | 2005        | PS2         | 1400                 |
| TI            | ALL             | 2006        | PS2         | 1500                 |
| TI            | ALL             | 2006        | Wii         | 500                  |
| TI            | ALL             | 2005        | XBox 360    | 1700                 |
| L             | Sydney          | ALL         | ALL         | 3400                 |
| L             | Melbourne       | ALL         | ALL         | 1700                 |
| T             | ALL             | 2005        | ALL         | 3100                 |
| T             | ALL             | 2006        | ALL         | 2000                 |
| I             | ALL             | ALL         | PS2         | 2900                 |
| I             | ALL             | ALL         | Wii         | 500                  |
| I             | ALL             | ALL         | XBox 360    | 1700                 |
|               | ALL             | ALL         | ALL         | 5100                 |

2. See below.

```

SELECT  L, T, I, SUM(M)
FROM    R
GROUP BY L, T, I
UNION ALL
SELECT  L, T, ALL, SUM(M)
FROM    R
GROUP BY L, T
UNION ALL
SELECT  L, ALL, I, SUM(M)
FROM    R
GROUP BY L, I
UNION ALL
SELECT  ALL, T, I, SUM(M)
FROM    R
GROUP BY T, I
UNION ALL
SELECT  L, ALL, ALL, SUM(M)
FROM    R
GROUP BY L
UNION ALL
SELECT  ALL, T, ALL, SUM(M)
FROM    R
GROUP BY T
UNION ALL
SELECT  ALL, ALL, I, SUM(M)
FROM    R
GROUP BY I
UNION ALL
SELECT  ALL, ALL, ALL, SUM(M)
FROM    R

```

3. The iceberg cube is

| <i>cuboid</i> | <i>Location</i> | <i>Time</i> | <i>Item</i> | <i>SUM(Quantity)</i> |
|---------------|-----------------|-------------|-------------|----------------------|
| LT            | Sydney          | 2006        | ALL         | 2000                 |
| LI            | Sydney          | ALL         | PS2         | 2900                 |
| L             | Sydney          | ALL         | ALL         | 3400                 |
| T             | ALL             | 2005        | ALL         | 3100                 |
| T             | ALL             | 2006        | ALL         | 2000                 |
| I             | ALL             | ALL         | PS2         | 2900                 |
|               | ALL             | ALL         | ALL         | 5100                 |

4. The mapping function we choose should satisfy the property that it is a one-to-one function (such that we can always recover the original value even after the mapping). The simplest form is  $h(L, T, I) = 12L + 4T + I$ . Hence,

| <i>Location</i> | <i>Time</i> | <i>Item</i> | <i>SUM(Quantity)</i> | <i>h(L, T, I)</i> |
|-----------------|-------------|-------------|----------------------|-------------------|
| 1               | 1           | 1           | 1400                 | 17                |
| 1               | 2           | 1           | 1500                 | 21                |
| 1               | 2           | 3           | 500                  | 23                |
| 2               | 1           | 2           | 1700                 | 30                |
| 1               | 1           | 0           | 1400                 | 16                |
| 1               | 2           | 0           | 2000                 | 20                |
| 2               | 1           | 0           | 1700                 | 28                |
| 1               | 0           | 1           | 2900                 | 13                |
| 1               | 0           | 3           | 500                  | 15                |
| 2               | 0           | 2           | 1700                 | 26                |
| 0               | 1           | 1           | 1400                 | 5                 |
| 0               | 2           | 1           | 1500                 | 9                 |
| 0               | 2           | 3           | 500                  | 11                |
| 0               | 1           | 2           | 1700                 | 6                 |
| 1               | 0           | 0           | 3400                 | 12                |
| 2               | 0           | 0           | 1700                 | 24                |
| 0               | 1           | 0           | 3100                 | 4                 |
| 0               | 2           | 0           | 2000                 | 8                 |
| 0               | 0           | 1           | 2900                 | 1                 |
| 0               | 0           | 3           | 500                  | 3                 |
| 0               | 0           | 2           | 1700                 | 2                 |
| 0               | 0           | 0           | 5100                 | 0                 |

So the final result is:

| <i>index</i> | <i>value</i> |
|--------------|--------------|
| 17           | 1400         |
| 21           | 1500         |
| 23           | 500          |
| 30           | 1700         |
| 16           | 1400         |
| 20           | 2000         |
| 28           | 1700         |
| 13           | 2900         |
| 15           | 500          |
| 26           | 1700         |
| 5            | 1400         |
| 9            | 1500         |
| 11           | 500          |
| 6            | 1700         |
| 12           | 3400         |
| 24           | 1700         |
| 4            | 3100         |
| 8            | 2000         |
| 1            | 2900         |
| 3            | 500          |
| 2            | 1700         |
| 0            | 5100         |

### Q3

1 For each chunk, we need to keep track of:

- the dimension that a new value is added (hence resulting in the current chunk)
- the size of all dimensions after the dimension value is added.

To enable fast mapping, for each value in each dimension, we also keep track of:

- the chunk ID for which this dimension value is associated with.

The following tables show the info kept.

| Chunk | Dimension | Dimension Sizes |
|-------|-----------|-----------------|
| 1     | 0         | $3 \times 2$    |
| 2     | $A$       | $4 \times 2$    |
| 3     | $B$       | $4 \times 3$    |
| 4     | $A$       | $5 \times 3$    |
| 5     | $A$       | $6 \times 3$    |
| 6     | $B$       | $6 \times 4$    |

| Dimension | Value | Chunk |
|-----------|-------|-------|
| $A$       | 0     | 1     |
| $A$       | 1     | 1     |
| $A$       | 2     | 1     |
| $A$       | 3     | 2     |
| $A$       | 4     | 4     |
| $A$       | 5     | 5     |
| $B$       | 0     | 1     |
| $B$       | 1     | 1     |
| $B$       | 2     | 3     |
| $B$       | 3     | 6     |

**2** We can then perform the mappings.

- A new chunk, 7, is created, of size  $1 \times 4$ . Two new entries are inserted to the above two tables, which are
  - $(7, A, 7 \times 4)$
  - $(A, 6, 7)$
- Given the logical address  $(1, 3)$ , we identify its chunk which is the larger chunk ID associated with dimension values. In this case, it is  $\max(1, 6) = 6$ . The chunk 6 is created when a new value is appended to Dimension  $B$ , hence, the offset within the chunk is based on its Dimension  $A$  value. Therefore, it is the 2nd entry in the chunk (as the index starts from 0).
- Given the physical address of the last entry (i.e., the 3rd entry) in Chunk 4, we first get Chunk 4's dimension sizes, which is  $5 \times 3$ , and we also know the new value is appended to Dimension  $A$ . Therefore, its logical address must have the new value, 4, associated with Dimension  $A$ , and its Dimension  $B$  value is the third value on Dimension  $B$ , which is 2.

**3**

- $O(1)$  for inserting expanding a dimension.
- $O(1)$  for  $g()$ , if the 2nd table is organized as a hash table mapping  $(Dimension, Value)$  to  $Chunk$ .
- $O(1)$  for  $g^{-1}()$ .